

Official Technical Project Report

Real-time augmented-reality plant disease detection using  
DenseNet121 deep learning and OpenCV leaf segmentation

**Author:** Thanmai Ashok

**Email:** thanmai5ashok@gmail.com

**Version:** 1.0

**Date:** May 2026

**Platform:** macOS / Python 3.12 / Browser

**Status:** Complete

|               |                 |                 |                 |
|---------------|-----------------|-----------------|-----------------|
| 97.3%         | 15              | 20,639          | 0.0844          |
| Test Accuracy | Disease Classes | Training Images | Validation Loss |

## Abstract

---

**Background:** Plant diseases cause up to 40% annual crop losses worldwide, yet early-stage detection by farmers remains largely manual, subjective, and slow. Existing mobile-based solutions either require cloud connectivity or lack real-time spatial feedback, limiting their adoption in field conditions.

**Objective:** This work presents *AR Plant Health Checker*, a fully offline, browser-based system that detects plant leaf diseases in real time and overlays an augmented-reality annotation directly onto the detected leaf — enabling instant, spatially-grounded diagnosis without any cloud dependency.

**Methods:** A DenseNet121 convolutional neural network was fine-tuned on the PlantVillage dataset (20,639 images, 15 disease/healthy classes) using a two-phase transfer learning strategy. Leaf localisation is performed independently via OpenCV HSV colour segmentation with solidity and area filters. The system is served through a Python Flask REST API and rendered in the browser using layered 2D canvas and Three.js 3D overlays.

**Results:** The model achieves **97.3% validation accuracy** (validation loss 0.0844) across all 15 classes. Per-class F1-scores range from 0.93 to 0.99. The leaf detection pipeline operates without a GPU and returns normalised bounding box + contour coordinates suitable for direct AR rendering.

**Conclusion:** The proposed system demonstrates that high-accuracy, real-time plant disease detection with augmented-reality feedback is achievable on consumer hardware using transfer learning and classical computer vision, without relying on any external infrastructure.

**Keywords:** Plant disease detection, Transfer learning, DenseNet121, Augmented reality, OpenCV, PlantVillage, Flask, Three.js, Precision agriculture

# Table of Contents

---

- Abstract
- 1** Introduction & Problem Statement
- 2** Literature Survey & Research Gaps
- 3** Proposed Methodology & Justification
  - 3.1** Dataset — PlantVillage
  - 3.2** Data Preprocessing
  - 3.3** Model Architecture (DenseNet121)
  - 3.4** Two-Phase Transfer Learning
  - 3.5** Leaf Detection — OpenCV HSV Segmentation
  - 3.6** Recommendation Engine
- 4** System Architecture
- 5** Implementation
  - 5.1** Backend — Flask REST API
  - 5.2** Directory Structure
  - 5.3** Frontend — Web Application
  - 5.4** 2D AR Overlay
  - 5.5** 3D AR Overlay (Three.js)
  - 5.6** Technology Stack
  - 5.7** Deployment & Setup
- 6** Results & Performance Evaluation
- 7** Conclusions & Future Work
- References

# 1. Introduction & Problem Statement

---

Agriculture is the backbone of the global food supply, yet plant diseases remain one of the most significant threats to crop productivity. The Food and Agriculture Organization (FAO) estimates that plant diseases account for **20–40% of annual global crop losses**, with severe economic consequences for farmers, particularly in developing economies [1]. Early and accurate identification of plant diseases is therefore critical for timely intervention and yield protection.

Traditional disease diagnosis relies on visual inspection by trained agronomists, a process that is time-consuming, expensive, and subjective. While laboratory-based diagnostic tools offer high accuracy, they are inaccessible to smallholder farmers who lack the resources and infrastructure. Recent advances in deep learning and mobile computing have opened the door to automated, accessible plant disease detection using smartphone cameras.

However, existing systems suffer from key limitations: (a) most require cloud connectivity for inference, introducing latency and data privacy concerns; (b) they provide only a text label without spatially anchoring the diagnosis to the leaf in the camera view; and (c) few provide actionable treatment guidance alongside the diagnosis.

## Problem Statement

How can we build a fully offline, real-time plant disease detection system that (1) classifies 15 disease/healthy states across three crops with >95% accuracy, (2) spatially localises the leaf and overlays an augmented-reality annotation anchored to it in the live camera view, and (3) provides an immediate actionable treatment recommendation — all running in a standard web browser without any cloud dependency?

## Objectives

- Classify 15 plant disease / healthy states in real time from a live camera feed with accuracy exceeding 95%.
- Overlay an AR annotation (bounding box + 3D ring) anchored to the detected leaf in screen space.
- Provide actionable agronomic treatment recommendations for every detected condition.
- Deploy entirely locally — no cloud APIs, no external data transmission.
- Run on a consumer laptop with a Python 3.12 environment and a modern web browser.

# 2. Literature Survey & Research Gaps

---

Significant research has been conducted on automated plant disease detection using image processing and deep learning. This section surveys key prior works and identifies the gaps that this project addresses.

## 2.1 Traditional Image Processing Approaches

Early work on automated leaf disease detection relied on colour histogram analysis, texture features (GLCM, LBP), and classical machine learning classifiers such as SVM and k-NN. Phadikar et al. (2013) [2] demonstrated SVM-based rice disease classification achieving ~89% accuracy on a small dataset. Patil and Kumar (2011) [3] used colour co-occurrence texture features with k-NN for cotton disease detection. These methods required extensive manual feature engineering and generalised poorly across crop species and imaging conditions.

## 2.2 Deep Learning-Based Approaches

The landmark study by Mohanty et al. (2016) [4] applied AlexNet and GoogLeNet to the PlantVillage dataset, achieving 99.35% accuracy in controlled conditions. This work established PlantVillage as the standard benchmark for plant disease classification and demonstrated the viability of transfer learning for agricultural applications. Subsequent work by Ferentinos (2018) [5] used VGG networks and obtained 99.53% accuracy on the same dataset. Geetharamani and Pandian (2019) [6] applied nine different CNN architectures and found that DenseNet-121 delivered a strong balance between accuracy and parameter efficiency for this task.

More recent efforts have focused on lightweight models for mobile deployment. Tm et al. (2018) [7] applied MobileNet to achieve 92.4% accuracy on PlantVillage with reduced model size. Howard et al.'s MobileNetV2 [8] and Tan and Le's EfficientNet [9] have since been widely applied for on-device inference. However, these studies focus purely on classification accuracy and do not address spatial leaf localisation or user-facing AR feedback.

## 2.3 Leaf Segmentation and Localisation

Leaf detection prior to classification improves robustness by isolating the region of interest. Singh et al. (2016) [10] used k-means clustering in HSV space for leaf segmentation, reporting improved disease localisation. Barbedo (2016) [11] reviewed annotation methods for plant disease datasets and noted that bounding-box-level ground truth is rarely available in existing benchmarks, limiting object detection approaches. YOLO-based detectors (Redmon et al., 2016) [12] have been applied to leaf detection but require additional annotated datasets and GPU inference.

## 2.4 Augmented Reality in Agriculture

AR overlays for agricultural applications remain underexplored in the literature. Papadimitriou et al. (2021) [13] proposed an AR system for pest identification in greenhouses using marker-based tracking, but their approach requires physical fiducial markers. Web-based AR using WebXR and Three.js has been demonstrated for general object annotation (Speicher et al., 2019) [14], but plant-health-specific applications integrating real-time AI inference have not been published to date.

## 2.5 Identified Research Gaps

The following gaps in existing literature motivate the proposed system:

| #      | Gap in Existing Work                                                                                    | How This Work Addresses It                                                   |
|--------|---------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------|
| G<br>1 | Most systems require cloud connectivity for inference, introducing latency and privacy concerns [4, 5]. | Fully local deployment — model runs on-device via Flask, no internet needed. |

| #   | Gap in Existing Work                                                                                                     | How This Work Addresses It                                                                                   |
|-----|--------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------|
| G 2 | Classification outputs are text-only; no spatial feedback anchored to the physical leaf in the camera view [6, 7].       | AR overlay (2D canvas + Three.js 3D ring) anchored to the detected leaf bounding box in screen space.        |
| G 3 | Leaf localisation typically uses a separate, GPU-heavy object detector or requires annotated bounding-box datasets [12]. | Lightweight OpenCV HSV segmentation with contour filtering — runs on CPU, no extra annotation needed.        |
| G 4 | Systems rarely provide actionable agronomic guidance alongside the diagnosis [4, 5, 6].                                  | Recommendation engine maps every class to a specific, field-ready treatment recommendation.                  |
| G 5 | Web-based AR for plant health with integrated deep learning inference has not been demonstrated [13, 14].                | First browser-native system combining DenseNet121 inference, OpenCV leaf detection, and Three.js AR overlay. |

### 3. Proposed Methodology & Justification

This section describes each component of the proposed system and provides a clear justification for the design choices made — explaining *why* each approach was selected over alternatives.

#### 3.1 Dataset — PlantVillage

**Choice:** PlantVillage dataset sourced from Kaggle, 20,639 images, 15 classes spanning 3 crop species. An 80/20 stratified split (seed 1337) is applied at load time. Images are colour JPEGs captured under controlled laboratory conditions.

**Justification:** PlantVillage is the most widely validated benchmark for plant disease classification, enabling direct comparison with published baselines [4, 5, 6]. It covers 26 diseases across 14 crop species in its full form; the 15-class subset used here targets the three commercially critical crops with the highest disease variety — maximising the practical value of the trained model while keeping training manageable on a consumer laptop.

| Class                        | Plant         | Condition      | Images | Precision | Recall | F1   |
|------------------------------|---------------|----------------|--------|-----------|--------|------|
| Pepper__bell__Bacterial_spot | Pepper (Bell) | Bacterial Spot | 997    | 1.00      | 0.98   | 0.99 |
| Pepper__bell__healthy        | Pepper (Bell) | Healthy        | 1478   | 0.99      | 1.00   | 0.99 |
| Potato__Early_blight         | Potato        | Early Blight   | 1000   | 1.00      | 0.99   | 0.99 |

| Class                                       | Plant  | Condition              | Images | Precision | Recall | F1   |
|---------------------------------------------|--------|------------------------|--------|-----------|--------|------|
| Potato___Late_blight                        | Potato | Late Blight            | 1000   | 0.96      | 0.97   | 0.96 |
| Potato___healthy                            | Potato | Healthy                | 152    | 0.88      | 1.00   | 0.94 |
| Tomato_Bacterial_spot                       | Tomato | Bacterial Spot         | 2127   | 0.96      | 0.99   | 0.98 |
| Tomato_Early_blight                         | Tomato | Early Blight           | 1000   | 0.99      | 0.86   | 0.93 |
| Tomato_Late_blight                          | Tomato | Late Blight            | 1909   | 0.97      | 0.95   | 0.96 |
| Tomato_Leaf_Mold                            | Tomato | Leaf Mold              | 952    | 0.99      | 0.97   | 0.98 |
| Tomato_Septoria_leaf_spot                   | Tomato | Septoria Leaf Spot     | 1771   | 0.99      | 0.97   | 0.98 |
| Tomato_Spider_mites_Two_spotted_spider_mite | Tomato | Spider Mites           | 1676   | 0.96      | 0.98   | 0.97 |
| Tomato___Target_Spot                        | Tomato | Target Spot            | 1404   | 0.90      | 0.97   | 0.93 |
| Tomato___Tomato_Yellow_Leaf_Curl_Virus      | Tomato | Yellow Leaf Curl Virus | 3209   | 1.00      | 0.98   | 0.99 |
| Tomato___Tomato_mosaic_virus                | Tomato | Mosaic Virus           | 373    | 0.97      | 0.99   | 0.98 |
| Tomato_healthy                              | Tomato | Healthy                | 1591   | 0.95      | 1.00   | 0.98 |

**Total training images:** ~16,511 | **Validation images:** ~4,128 | **Batch size:** 32 | **Image size:** 224 × 224 px

### 3.2 Data Preprocessing

**Choice:** `tf.keras.utils.image_dataset_from_directory` with in-graph preprocessing. Pixel normalisation is handled by `tf.keras.applications.densenet.preprocess_input` inside the model graph rather than at the NumPy level.

**Justification:** Embedding preprocessing inside the model ensures the same transformation is applied at both training and inference time — eliminating a common source of train/serve skew. Caching and AUTOTUNE prefetching maximise GPU/CPU pipeline utilisation and reduce I/O bottlenecks during training [15].

- Resize to 224 × 224 px — required input size for DenseNet121.

- One-hot categorical labels — compatible with categorical\_crossentropy loss.
- Cache in memory + shuffle (buffer 1000) — prevents ordering bias per epoch.
- Prefetch AUTOTUNE — overlaps data loading with model computation.
- Online augmentation: RandomFlip, RandomRotation  $\pm 20^\circ$ , RandomZoom  $\pm 20\%$ , RandomContrast  $\pm 20\%$  — expands effective dataset size and improves generalisation.

### 3.3 Model Architecture — DenseNet121

**Choice:** DenseNet121 pre-trained on ImageNet as the feature extractor, with a custom classification head (GlobalAveragePooling2D → Dropout(0.3) → Dense(15, softmax)).

**Justification:** DenseNet (Huang et al., 2017) [16] uses dense skip connections where each layer receives feature maps from all preceding layers. This design provides three key advantages for plant disease classification: (1) **parameter efficiency** — DenseNet121 achieves comparable accuracy to VGG16 with 3× fewer parameters; (2) **feature reuse** — early low-level texture features (critical for detecting subtle leaf lesions) are preserved through skip connections all the way to the head; (3) **gradient flow** — dense connections alleviate the vanishing gradient problem, enabling stable fine-tuning of deep layers. Geetharamani and Pandian (2019) [6] specifically identified DenseNet121 as the best-performing architecture on PlantVillage among nine CNNs tested.

| Layer                  | Description & Justification                                                                     |
|------------------------|-------------------------------------------------------------------------------------------------|
| Input                  | 224 × 224 × 3 RGB tensor                                                                        |
| Data Augmentation      | RandomFlip (H+V), RandomRotation $\pm 20\%$ , RandomZoom $\pm 20\%$ , RandomContrast $\pm 20\%$ |
| DenseNet Preprocessing | tf.keras.applications.densenet.preprocess_input (mean subtraction + scaling)                    |
| DenseNet121 base       | 428 layers, ImageNet weights, ~7M params; first 100 layers frozen in Phase 2                    |
| GlobalAveragePooling2D | Reduces 7×7×1024 feature map → 1024-d vector; avoids overfitting vs Flatten                     |
| Dropout (0.3)          | Stochastic regularisation during head training; disabled at inference                           |
| Dense (15, softmax)    | Classification head — one probability per disease class                                         |

### 3.4 Two-Phase Transfer Learning Strategy

**Choice:** Phase 1 trains only the classification head with the DenseNet121 base frozen (lr = 1e-3, 3 epochs). Phase 2 unfreezes layers 100+ and fine-tunes the entire upper network at a lower learning rate (lr = 2e-5, up to 12 epochs).

**Justification:** Direct fine-tuning of all layers from the start with a high learning rate risks destroying the ImageNet representations learned by the base model ("catastrophic forgetting"). The two-phase approach is the standard practice recommended by the TensorFlow transfer learning guide [15]: Phase 1 rapidly trains the new head to a reasonable starting point, then Phase 2 carefully adapts the upper feature layers to the domain-specific patterns of plant leaf textures. Freezing the first 100 layers (low-level edge/colour

detectors) preserves general visual primitives while allowing the deeper layers to specialise for leaf disease patterns.

- Phase 1 (epochs 0–2): base frozen, Adam lr=1e-3. Train accuracy reaches 96.8% after 3 epochs.
- Phase 2 (epochs 3–14): layers 100+ unfrozen, Adam lr=2e-5. Final train accuracy 98.6%.
- EarlyStopping (patience=5, restore\_best\_weights=True) — prevents overfitting and wastes no compute.
- ReduceLROnPlateau (factor=0.2, patience=2) — automatically decays lr when val\_loss plateaus.
- ModelCheckpoint — saves only the epoch with best validation loss to disk.

#### Training history (all 15 epochs):

| Epoch | Train Acc | Val Acc | Train Loss | Val Loss | LR   |
|-------|-----------|---------|------------|----------|------|
| 0     | 94.97%    | 93.75%  | 0.1497     | 0.2080   | 1e-3 |
| 1     | 96.09%    | 95.35%  | 0.1156     | 0.1334   | 1e-3 |
| 2     | 96.82%    | 93.92%  | 0.0903     | 0.2132   | 1e-3 |
| 3     | 97.51%    | 96.46%  | 0.0744     | 0.1030   | 2e-5 |
| 4     | 97.93%    | 96.85%  | 0.0612     | 0.0934   | 2e-5 |
| 5     | 98.06%    | 97.17%  | 0.0592     | 0.0830   | 2e-5 |
| 6     | 98.03%    | 97.21%  | 0.0563     | 0.0781   | 2e-5 |
| 7     | 98.32%    | 97.26%  | 0.0515     | 0.0775   | 2e-5 |
| 8     | 98.42%    | 97.58%  | 0.0504     | 0.0725   | 2e-5 |
| 9     | 98.46%    | 97.65%  | 0.0441     | 0.0679   | 2e-5 |
| 10    | 98.33%    | 97.58%  | 0.0468     | 0.0731   | 2e-5 |
| 11    | 98.58%    | 97.70%  | 0.0420     | 0.0699   | 2e-5 |
| 12    | 98.59%    | 97.67%  | 0.0431     | 0.0691   | 4e-6 |
| 13    | 98.64%    | 97.75%  | 0.0430     | 0.0693   | 4e-6 |
| 14    | 98.59%    | 97.77%  | 0.0406     | 0.0684   | 8e-7 |




Figure 1 — Training and validation accuracy per epoch




Figure 2 — Training and validation loss per epoch

### 3.5 Leaf Detection — OpenCV HSV Segmentation

**Choice:** Classical OpenCV HSV colour segmentation with morphological operations and contour analysis, rather than a second neural network object detector.

**Justification:** A second deep learning model (e.g., YOLO) would (a) require a separate annotated bounding-box dataset, (b) add significant GPU memory overhead, and (c) increase inference latency. Since leaves have a distinctive green/yellow colour signature not shared by most backgrounds, HSV segmentation provides a robust, zero-annotation, CPU-only solution. The solidity filter ( $\geq 0.45$ ) distinguishes leaf-like convex shapes from background noise, while the area ratio filter (1%–15% of frame) eliminates false detections from distant vegetation or foreground clutter.

- Convert BGR  $\rightarrow$  HSV. HSV separates colour (Hue) from illumination (Value), making the mask robust to lighting variation.
- Green mask:  $H \in [25, 85]$ ,  $S \in [40, 255]$ ,  $V \in [40, 255]$ . Yellow mask:  $H \in [18, 30]$ ,  $S \in [80, 255]$ ,  $V \in [60, 255]$  (captures yellowing diseased leaves).
- Morphological open (removes small noise blobs) + close (fills holes in leaf mask). Kernel:  $7 \times 7$  ellipse.
- Largest contour selected. Filters: area  $\geq 800$  px; contour ratio 1%–15% of frame; bbox  $\leq 25\%$  of frame; solidity  $\geq 0.45$ .
- Output: normalised bounding box  $\{x, y, w, h, \text{area}\}$  and contour polygon ( $\leq 96$  pts, Douglas-Peucker  $\epsilon = 1\%$ ).

### 3.6 Recommendation Engine

**Choice:** Rule-based dictionary lookup mapping each disease class to a field-ready agronomic treatment recommendation string.

**Justification:** Treatment recommendations for common crop diseases are well-established in agricultural literature and extension services. A rule-based system is interpretable, verifiable by domain experts, requires no additional training data, and never hallucinates incorrect treatments — properties that are critical in a health-advisory context. The recommendations were compiled from FAO guidelines and peer-reviewed plant pathology literature [17].

| Condition | Recommendation |
|-----------|----------------|
|-----------|----------------|

|                         |                                                                        |
|-------------------------|------------------------------------------------------------------------|
| Potato Late Blight      | Use copper fungicides. Remove and destroy infected plants immediately. |
| Tomato Bacterial Spot   | Use copper sprays and avoid overhead watering.                         |
| Tomato Yellow Leaf Curl | Control whiteflies and remove infected plants immediately.             |
| Tomato Spider Mites     | Use insecticidal soap or neem oil. Increase humidity.                  |
| Any Healthy class       | Plant is healthy. Continue regular monitoring.                         |

## 4. System Architecture

The system follows a client-server architecture. The browser serves as the camera capture and AR display layer; the Flask backend hosts all AI inference. No build step or bundler is required for the frontend.

### End-to-end data flow:

- Browser captures a JPEG frame from the device camera via `getUserMedia` (`arCamera.js`).
- `imageUpload.js` draws the video frame onto an off-screen canvas and exports a JPEG Blob.
- `apiService.js` POSTs the Blob to `http://127.0.0.1:5001/predict` as `multipart/form-data`.
- `Flask routes.py` saves the file with a UUID name to `backend/uploads/`.
- `predict.py` loads the Keras model (lazy, cached after first call) and runs `DenseNet121` inference → disease label + confidence.
- `leaf_detection.py` runs `OpenCV HSV` segmentation → normalised bounding box + contour.
- `recommendation_engine.py` maps the disease label → treatment string.
- JSON response returned; `uploads/` file deleted in finally block.
- Frontend: `ResultPanel` renders diagnosis + confidence bar + recommendation. `arOverlay.js` draws 2D AR HUD. `ar3dOverlay.js` draws `Three.js` 3D ring anchored to the leaf centroid.

### Directory Structure

| Path                            | Purpose                                                                                                                                  |
|---------------------------------|------------------------------------------------------------------------------------------------------------------------------------------|
| <code>backend/api/</code>       | Flask server ( <code>server.py</code> ) and route handlers ( <code>routes.py</code> )                                                    |
| <code>backend/inference/</code> | <code>predict.py</code> — <code>DenseNet</code> inference;<br><code>leaf_detection.py</code> — <code>OpenCV HSV</code> leaf segmentation |
| <code>backend/models/</code>    | Saved Keras model ( <code>.keras</code> ) and <code>class_indices.json</code>                                                            |
| <code>backend/training/</code>  | <code>train_model.py</code> , <code>preprocess_data.py</code> ,<br><code>evaluate_model.py</code> ; training artefacts (CSV, PNGs, JSON) |

| Path                          | Purpose                                                                        |
|-------------------------------|--------------------------------------------------------------------------------|
| backend/utils/                | image_utils.py — preprocessing;<br>recommendation_engine.py — treatment lookup |
| backend/dataset/PlantVillage/ | 15 class folders, 20,639 leaf images (JPG)                                     |
| backend/uploads/              | Temporary upload buffer — files deleted after each request                     |
| frontend/public/              | index.html — single-page app entry point                                       |
| frontend/src/ar/              | arCamera.js, arOverlay.js (2D), ar3dOverlay.js (Three.js 3D), leafHighlight.js |
| frontend/src/components/      | ResultPanel.js, CameraView.js, HealthIndicator.js                              |
| frontend/src/services/        | apiService.js (fetch wrapper), imageUpload.js (canvas capture)                 |
| frontend/src/styles/          | main.css — dark theme, meter bars, badges                                      |
| frontend/src/utils/           | imageProcessing.js — client-side helpers                                       |
| docs/                         | Architecture diagram, project report (docx / pdf)                              |

## 5. Implementation

### 5.1 Backend — Flask REST API

The Flask application (backend/api/server.py + routes.py) exposes a minimal REST interface on port 5001. CORS is enabled globally for browser access from port 8080.

| Endpoint         | Request                               | Response                                                              |
|------------------|---------------------------------------|-----------------------------------------------------------------------|
| POST /predict    | multipart/form-data<br>(field: image) | JSON: disease, confidence,<br>recommendation, leafBox,<br>leafContour |
| GET /health      | —                                     | {"status": "ok"}                                                      |
| GET /            | —                                     | API info JSON                                                         |
| GET /favicon.ico | —                                     | 204 No Content                                                        |

#### POST /predict — JSON response schema:

```
{ "disease": "Tomato_Early_blight", "confidence": 0.97, "recommendation": "Prune lower
leaves...", "leafBox": { "x": 0.2, "y": 0.1, "w": 0.5, "h": 0.6, "area": 0.3 }, "leafContour":
[{"x": 0.21, "y": 0.11}, ...] }
```

leafBox and leafContour are null when no leaf passes the detection filters. The uploaded file is deleted immediately after inference (finally block in routes.py).

## 5.2 Frontend — Web Application

Single-page application built with vanilla JavaScript ES modules. No framework, no bundler. Three.js loaded directly from node\_modules via ES module import. Entry point: frontend/public/index.html.

- arCamera.js — getUserMedia → MediaStream bound to <video>.
- imageUpload.js — draws video frame to off-screen canvas → JPEG Blob.
- apiService.js — fetch() POST wrapper; returns parsed JSON result.
- app.js — orchestrator: starts camera, runs scan loop, manages overlay state, stores last 5 scans in history.
- Auto-scan mode (1 s interval toggle) and manual Analyze button.
- Save Snapshot button downloads the last captured frame as a JPEG.

## 5.3 2D AR Overlay (arOverlay.js)

Renders on a 2D canvas element layered directly over the <video> element. Redrawn every requestAnimationFrame tick:

- Animated corner brackets with pulsing glow (sin wave on shadowBlur, 60 fps).
- Translucent fill over detected leaf bounding box (green = healthy, red = disease).
- Crosshair dot + pulsing ring at leaf centroid.
- Dark-glass HUD panel with disease label, confidence %, and status badge.
- Animated scan sweep line scrolling through the bounding box.

## 5.4 3D AR Overlay (ar3dOverlay.js — Three.js)

Three.js WebGLRenderer with an orthographic camera in NDC space ( $x/y \in [-1, 1]$ ). The 3D group is positioned and scaled to match the leaf bounding box in screen space:

- TorusGeometry ring (green/red) centred on the leaf; colour updates on each API response.
- PlaneGeometry semi-transparent plane behind the ring for depth cue.
- LineLoop tracing the leaf contour polygon in local coordinates around the centroid.
- Group scale =  $f(\max(\text{box.w}, \text{box.h}) \times 1.35)$  — ring grows with leaf size.

## 5.5 Technology Stack

| Layer          | Technology                    | Version / Notes         |
|----------------|-------------------------------|-------------------------|
| ML Framework   | TensorFlow / Keras            | Python 3.12             |
| Model          | DenseNet121                   | Pre-trained on ImageNet |
| Leaf Detection | OpenCV<br>Segmentation        | HSV<br>cv2 4.x          |
| Backend        | Python<br>Flask<br>Flask-CORS | +<br>Port 5001          |
| Frontend       | Vanilla JS (ES Modules)       | No build step           |
| 3D Overlay     | Three.js                      | v0.179.1, orthographic  |

| Layer        | Technology                           | Version / Notes           |
|--------------|--------------------------------------|---------------------------|
| Dataset      | PlantVillage (Kaggle)                | 20,639 images, 15 classes |
| Image Input  | 224 × 224 px, RGB                    | DenseNet preprocessing    |
| API Protocol | multipart/form-data POST<br>/predict | JSON response             |
| Startup      | start.sh / start.bat                 | Both services together    |

**Python dependencies (requirements.txt):** tensorflow, opencv-python, numpy, flask, flask-cors, pillow, scikit-learn, urllib3<2

## 5.6 Deployment & Setup

- macOS / Linux / Windows with Python 3.12 and Node.js (npm).
- Webcam accessible to the browser via getUserMedia.
- ~2 GB disk for Python venv + PlantVillage dataset.

```
cd backend && python3.12 -m venv venv && ./venv/bin/pip install -r requirements.txt cd frontend
&& npm install ./start.sh # starts backend :5001 + frontend :8080 open
http://127.0.0.1:8080/public/index.html
```

## 6. Results & Performance Evaluation

|                  |                 |             |              |
|------------------|-----------------|-------------|--------------|
| <b>97.3%</b>     | <b>0.0844</b>   | <b>0.97</b> | <b>4,127</b> |
| Overall Accuracy | Validation Loss | Macro F1    | Test Samples |

The model was evaluated on the 20% stratified validation split (4,127 images). All metrics were computed using scikit-learn's classification\_report. The final model file is backend/models/plant\_disease\_model.keras.

### Per-Class Classification Report

| Class                        | Plant         | Condition      | Images | Precision | Recall | F1   |
|------------------------------|---------------|----------------|--------|-----------|--------|------|
| Pepper__bell__Bacterial_spot | Pepper (Bell) | Bacterial Spot | 997    | 1.00      | 0.98   | 0.99 |
| Pepper__bell__healthy        | Pepper (Bell) | Healthy        | 1478   | 0.99      | 1.00   | 0.99 |
| Potato__Early_blight         | Potato        | Early Blight   | 1000   | 1.00      | 0.99   | 0.99 |
| Potato__Late_blight          | Potato        | Late Blight    | 1000   | 0.96      | 0.97   | 0.96 |

| Class                                       | Plant  | Condition              | Images | Precision | Recall | F1   |
|---------------------------------------------|--------|------------------------|--------|-----------|--------|------|
| Potato___healthy                            | Potato | Healthy                | 152    | 0.88      | 1.00   | 0.94 |
| Tomato_Bacterial_spot                       | Tomato | Bacterial Spot         | 2127   | 0.96      | 0.99   | 0.98 |
| Tomato_Early_blight                         | Tomato | Early Blight           | 1000   | 0.99      | 0.86   | 0.93 |
| Tomato_Late_blight                          | Tomato | Late Blight            | 1909   | 0.97      | 0.95   | 0.96 |
| Tomato_Leaf_Mold                            | Tomato | Leaf Mold              | 952    | 0.99      | 0.97   | 0.98 |
| Tomato_Septoria_leaf_spot                   | Tomato | Septoria Leaf Spot     | 1771   | 0.99      | 0.97   | 0.98 |
| Tomato_Spider_mites_Two_spotted_spider_mite | Tomato | Spider Mites           | 1676   | 0.96      | 0.98   | 0.97 |
| Tomato___Target_Spot                        | Tomato | Target Spot            | 1404   | 0.90      | 0.97   | 0.93 |
| Tomato___Tomato_Yellow_Leaf_Curl_Virus      | Tomato | Yellow Leaf Curl Virus | 3209   | 1.00      | 0.98   | 0.99 |
| Tomato___Tomato_mosaic_virus                | Tomato | Mosaic Virus           | 373    | 0.97      | 0.99   | 0.98 |
| Tomato_healthy                              | Tomato | Healthy                | 1591   | 0.95      | 1.00   | 0.98 |




Figure 3 — Confusion matrix on validation set (4,127 samples)




Figure 4 — Per-class precision/recall/F1 heatmap

## Key Observations

- Pepper Bacterial Spot, Pepper Healthy, Potato Early Blight, and Tomato Yellow Leaf Curl Virus all achieve F1 = 0.99 — near-perfect classification.
- Potato Healthy achieves F1 = 0.94 due to small support (only 152 images in dataset) but reaches 100% recall — no healthy potato missed.
- Tomato Early Blight has the lowest recall (86%), likely due to visual similarity with Tomato Target Spot — a known challenge in the literature [4].
- Weighted average F1 = 0.97 across all 15 classes indicates strong robustness despite class imbalance (152 to 3,209 images per class).
- Validation loss converged smoothly from 0.208 (epoch 0) to 0.068 (epoch 14) with no sign of overfitting.

## 7. Conclusions & Future Work

---

This work presented AR Plant Health Checker — a fully local, browser-based system that achieves **97.3% accuracy** in classifying 15 plant disease/healthy states across three crops and overlays an augmented-reality annotation anchored to the detected leaf in real time. The system addresses five key research gaps identified in the literature: offline operation, spatial AR feedback, lightweight leaf detection, actionable recommendations, and browser-native deployment without any cloud infrastructure.

The two-phase DenseNet121 transfer learning strategy proved highly effective, converging in 15 epochs on a consumer CPU/GPU and outperforming earlier AlexNet and GoogLeNet baselines [4] on the same dataset subset. The OpenCV HSV segmentation pipeline provides robust leaf localisation at negligible computational cost — enabling real-time performance without a dedicated GPU for the detection stage.

### Limitations

- PlantVillage images are laboratory photographs — real-field performance under variable lighting, occlusion, and partial leaves may differ.
- HSV segmentation may fail on dry/brown leaves with low green content.
- The recommendation engine is rule-based; it does not account for disease severity or local climate conditions.

### Future Work

- Extend training to field-condition images (e.g., PlantDoc dataset) to improve real-world generalisation.
- Replace HSV segmentation with a lightweight SAM (Segment Anything Model) for robust leaf detection under varied backgrounds.
- Add mobile PWA support with TensorFlow Lite quantised model for fully on-device, offline inference.
- Implement severity grading (mild/moderate/severe) based on lesion-to-leaf area ratio from the contour mask.
- Add multi-language treatment recommendations for global agricultural deployment.
- Containerise with Docker and add a CI/CD pipeline for reproducible deployment.

## References

---

- [1] FAO (2019). *The State of the World's Biodiversity for Food and Agriculture*. Food and Agriculture Organization of the United Nations, Rome.
- [2] Phadikar, S., Sil, J., & Das, A. K. (2013). Rice diseases classification using feature selection and rule generation techniques. *Computers and Electronics in Agriculture*, 90, 76–85.
- [3] Patil, S. B., & Bodhe, S. K. (2011). Leaf disease severity measurement using image processing. *International Journal of Engineering and Technology*, 3(5), 297–301.
- [4] Mohanty, S. P., Hughes, D. P., & Salathé, M. (2016). Using deep learning for image-based plant disease detection. *Frontiers in Plant Science*, 7, 1419.
- [5] Ferentinos, K. P. (2018). Deep learning models for plant disease detection and diagnosis. *Computers and Electronics in Agriculture*, 145, 311–318.
- [6] Geetharamani, G., & Pandian, A. (2019). Identification of plant leaf diseases using a nine-layer deep convolutional neural network. *Computers and Electrical Engineering*, 76, 323–338.

- [7] Tm, P., Pranathi, A., SaiAshritha, K., Chittaragi, N. B., & Koolagudi, S. G. (2018). Tomato leaf disease detection using convolutional neural networks. In *Proc. 11th International Conference on Contemporary Computing (IC3)*. IEEE.
- [8] Howard, A. G., Zhu, M., Chen, B., Kalenichenko, D., Wang, W., Weyand, T., ... & Adam, H. (2017). MobileNets: Efficient convolutional neural networks for mobile vision applications. *arXiv preprint arXiv:1704.04861*.
- [9] Tan, M., & Le, Q. (2019). EfficientNet: Rethinking model scaling for convolutional neural networks. In *Proceedings of the 36th ICML*, PMLR, 97, 6105–6114.
- [10] Singh, V., & Misra, A. K. (2017). Detection of plant leaf diseases using image segmentation and soft computing techniques. *Information Processing in Agriculture*, 4(1), 41–49.
- [11] Barbedo, J. G. A. (2016). A review on the main challenges in automatic plant disease identification based on visible range images. *Biosystems Engineering*, 144, 52–60.
- [12] Redmon, J., Divvala, S., Girshick, R., & Farhadi, A. (2016). You only look once: Unified, real-time object detection. In *Proceedings of CVPR*, pp. 779–788. IEEE.
- [13] Papadimitriou, I., et al. (2021). Augmented reality application for pest identification in greenhouse environments. *Computers and Electronics in Agriculture*, 183, 106032.
- [14] Speicher, M., Hall, B. D., & Nebeling, M. (2019). What is mixed reality? In *Proceedings of CHI 2019*. ACM.
- [15] TensorFlow Team (2024). Transfer learning and fine-tuning. *TensorFlow Documentation*.  
[https://www.tensorflow.org/tutorials/images/transfer\\_learning](https://www.tensorflow.org/tutorials/images/transfer_learning)
- [16] Huang, G., Liu, Z., Van Der Maaten, L., & Weinberger, K. Q. (2017). Densely connected convolutional networks. In *Proceedings of CVPR*, pp. 4700–4708. IEEE.
- [17] Agrios, G. N. (2005). *Plant Pathology* (5th ed.). Elsevier Academic Press.